

BEE: A Scientific Application Workflow Engine

Manage Multi-step Simulations on
HPC & Cloud Platforms

Patricia Grubel, Quincy Wofford,
Rusty Davis



WoWoHa 2020 Summer Webinar
June 26, 2020

LA-UR-20-24727

Outline

- BEE Overview
- Charliecloud – Lightweight Container Runtime
- BEE Components & Internals
- Neo4j Graph Database
 - Managing Workflows
 - Archiving for recoverability, re-running, cloning
- HPC Scientific Workflows
- Example - CLAMR
- BEEstart – starts up the server components

BEE: A Scientific Application Workflow Engine

- Docker Image Support
 - Configurable support for HPC - Charliecloud, Shifter and Singularity Container Runtimes
- Supports Multiple HPC Clusters
 - Enables high resource usability
- Designed for HPC simulations
- Standard Workflow Specification: Common Workflow Language
- Automation
 - Platform-related setup, configuration and launching
 - Avoid learning arcane technical details of HPC resource managers
- No privileged access required
 - Any user can use on HPC platform of their choice
- Reproducible
 - Complex scientific workflows can be archived, share metadata, re-run
- Contact

bee-dev@lanl.gov

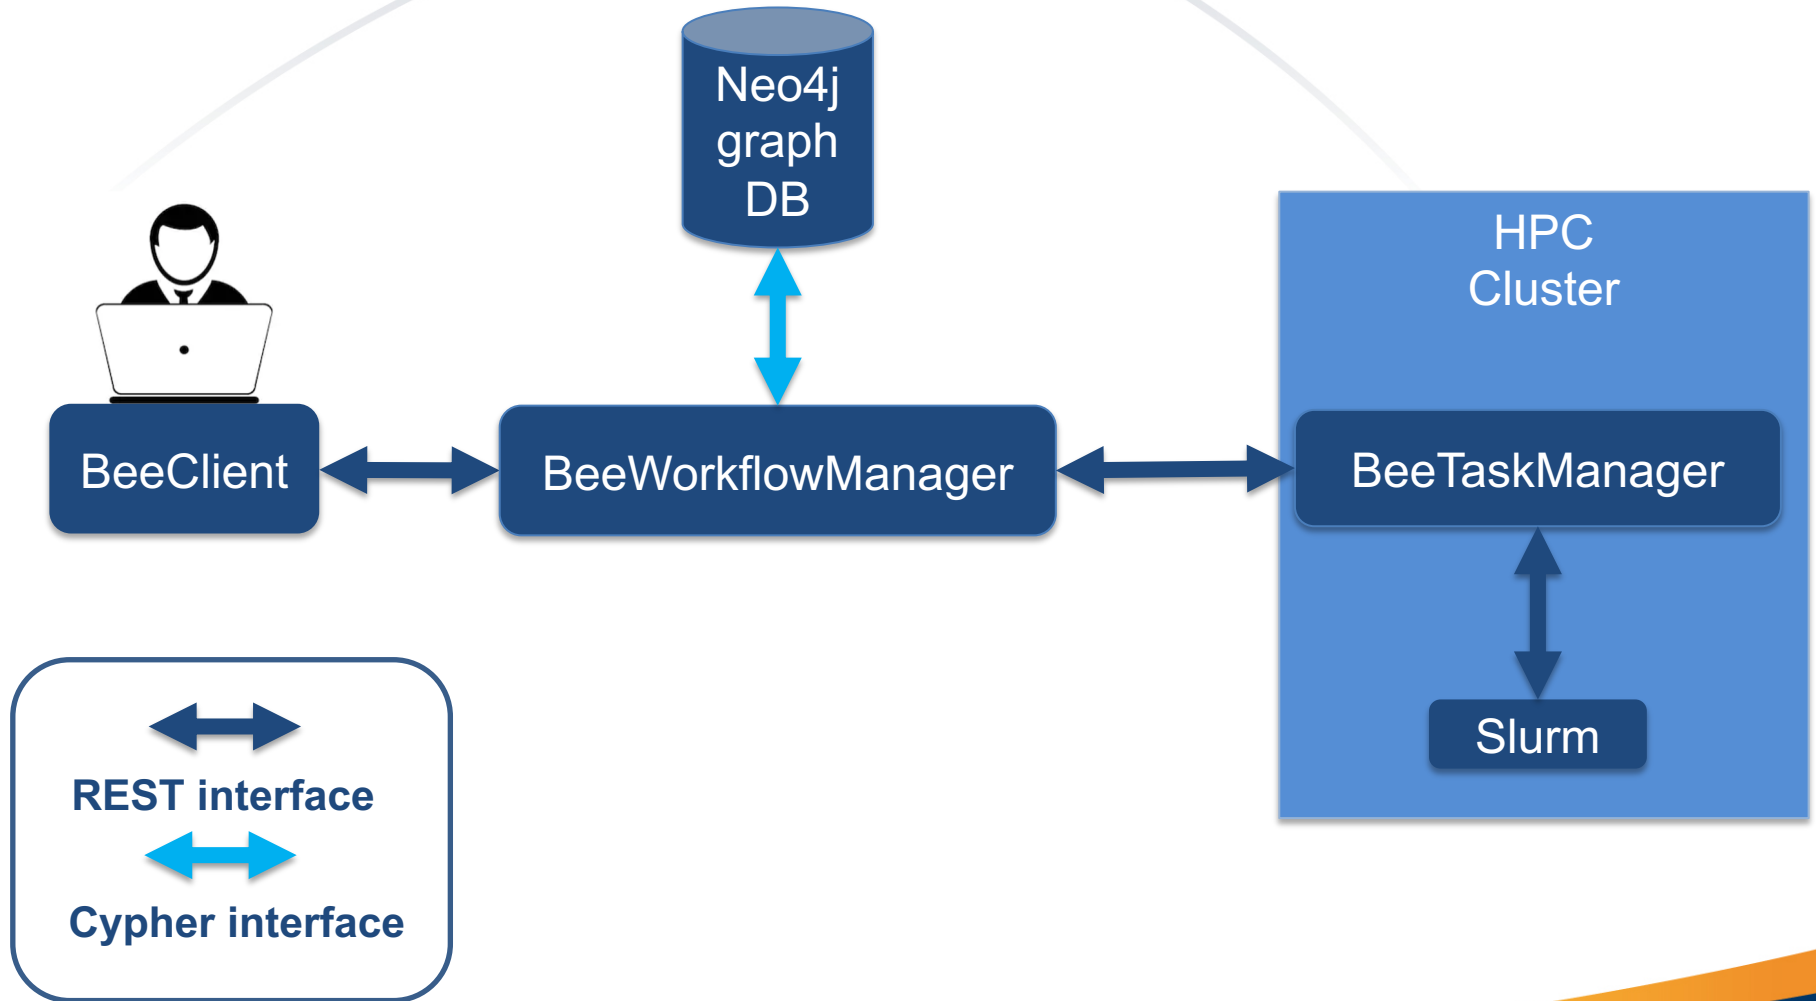


Charliecloud: HPC Container Runtime

- Lightweight
 - No complex installation
 - No configuration required
 - ~1300 Lines of Code (C & shell scripts)
- Unprivileged
 - Any user can run it
 - No HPC center admin intervention required
- Linux kernel support
 - All major distributions have kernel support for user namespaces
 - Some require a tweak (e.g. `/proc/sys/user/max_user_namespaces`)
- Hardware agnostic
 - Tested and supported on X86_64, ARM64, PowerPC
- Available today
 - <https://github.com/hpc/charliecloud>
 - Packages for all major Linux distributions
 - OpenHPC release



BEE: Components for HPC



BEE: Internals

- Python 3
 - Portable across Linux, OS X, Windows
 - Modular Design – Abstract Classes for Major Components
 - Support for using different Graph Database
 - Support for multiple Container Runtimes
 - Support for Multiple Workload Managers (Slurm, LSF, PBS, Torque/Moab...)
- REST and YAML
 - Easy to enhance and extend
- Common Workflow Language
 - Open standard, many tools already exist
 - Python for expressions instead of JavaScript
 - BEE extensions for better HPC support
 - Automatic Setup of HPC Requirements
 - HPC Container Runtimes, Charliecloud, Singularity etc.
- Neo4j - Graph database
 - Manage workflow and metadata
 - DAG allows workflow execution optimize
 - Archive workflow & artifacts



Neo4j - Graph Database


- Manage Workflow
 - Build Workflow DAG
 - Visualize DAG
 - Metadata during run – task state, job id
- Archive Workflow
 - Workflow metadata - what cluster, cluster job ids, times (submit, start, compute)
 - Provenance - ability to archive the workflow
 - Captures container UUID, input decks, run commands, checkpoint file location
 - Rerun workflows
 - Clone workflows - copy, reset data and go
 - Resiliency – true state in BEE is in the database
 - Recovery from outages
 - Component restart - components designed to continue using database metadata
 - Checkpoint / Restart

HPC Scientific Workflows - LANL

- Multi-Physics Simulation
 - Computations can last months
 - Checkpoint / Restart Capability is paramount
 - Repeatability
- Biology
 - Common Workflow Language
 - BEE will facilitate using HPC resources and containers
- Physics - Large Scale Parametric Studies
 - 1000's of Simulations
 - Diverging solutions can cause unneeded billing cycles
 - Facilitate ability to cancel, change input/parameters, then restart

CLAMR Workflow Example

- CLAMR
 - A LANL mini-app
 - Simulates shallow water equations
 - Performs hydrodynamic cell-based adaptive mesh refinement (AMR)
 - Intended as a testbed for hybrid algorithm development using MPI and OpenCL



CLAMR Visualization

CLAMR Workflow CWL

```
cwlVersion: v1.0

requirements:
  CommandLineRequirement: { }

id: clamr-flow
label: clamr-flow

steps:
  clamr:
    in:
      ...
    run: /clamr/CLAMR-master/clamr_cpunonly ...

  hints: DockerRequirement
    dockerImageID: ".../clamr.tar.gz"

  out: [ graphics_output/ ]

ffmpeg :
  clamr:
    in:
      ...
    run: / ffmpeg -f image2 -i graphics_output/ ...
    out: [ CLAMR_movie.mp4 ]
```

CLAMR Job Scripts

```
#!/bin/bash
#SBATCH
module load charliecloud

mkdir -p /tmp/$USER

ch-tar2dir /usr/projects/beedev/clamr/clamr-
toss.tar.gz /tmp/$USER

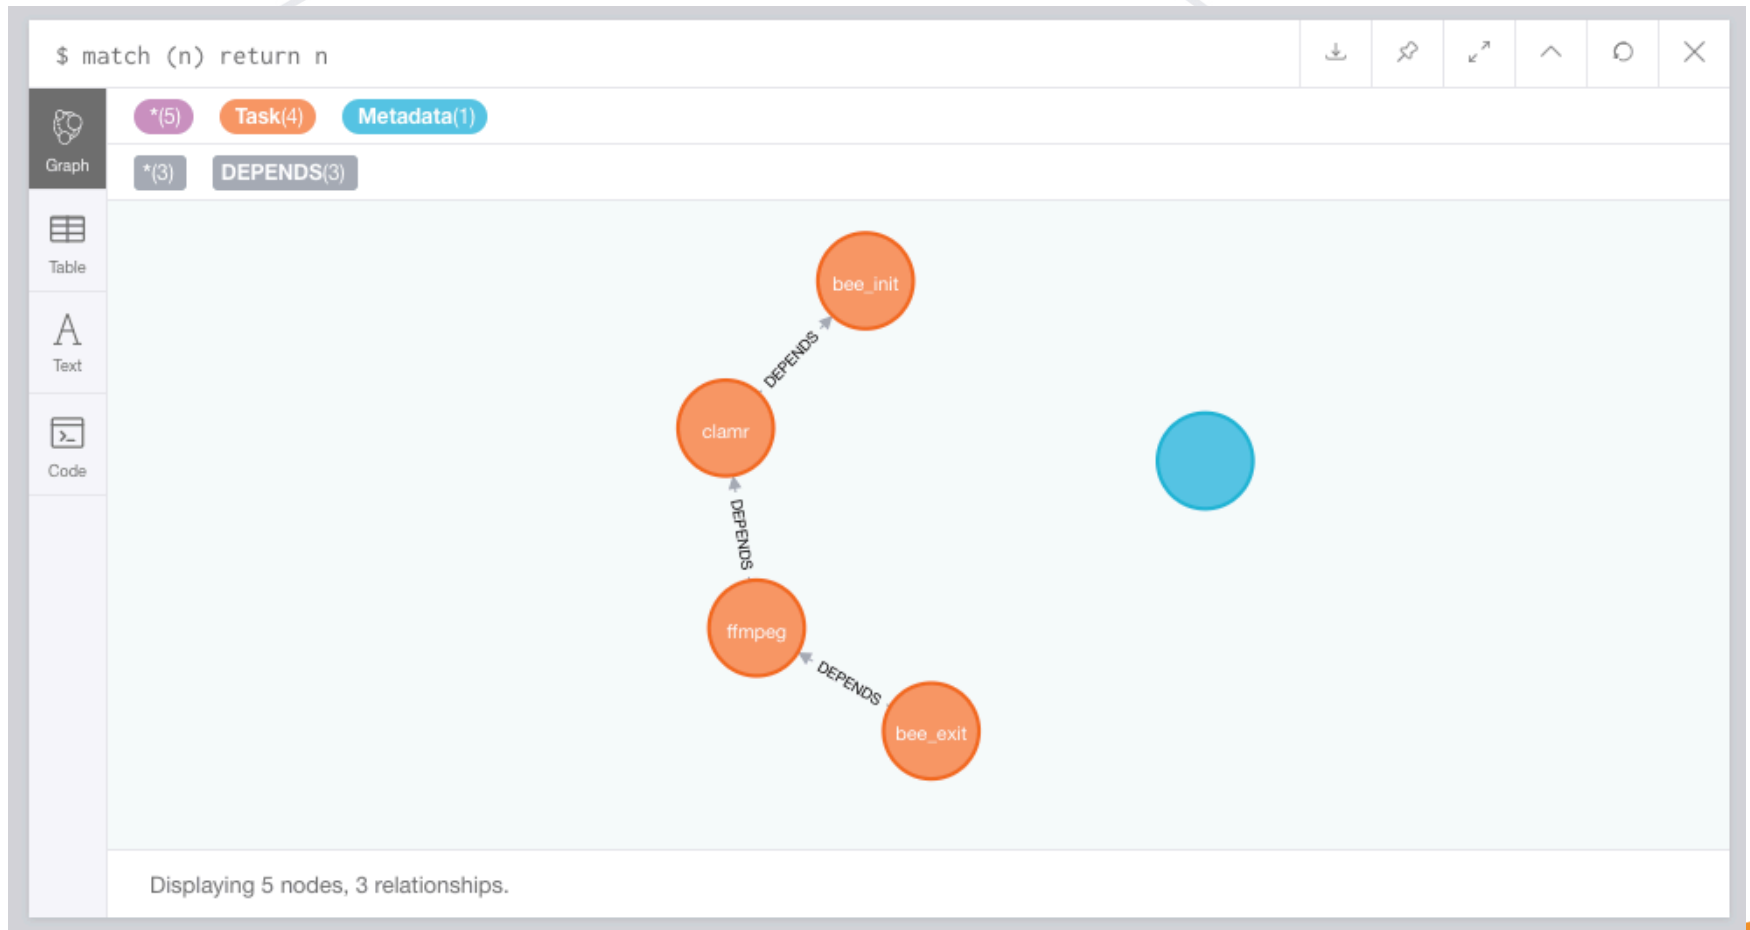
ch-run /tmp/$USER/clamr-toss -b $PWD -c
/mnt/0 -- /clamr/CLAMR-
master/clamr_cpuonly -n 32 -l 3 -t 5000 -i 10 -
g 100 -G png rm -rf /tmp/$USER/clamr-toss
```

CLAMR Job Step

```
#!/bin/bash
#SBATCH
ffmpeg -f image2 -i
graphics_output/graph%05d.png -r 12 -s
800x800 -pix_fmt yuv420p CLAMR_movie.mp4
```

ffmpeg Job Step

CLAMR neo4j Workflow



BEE Configuration File

- Orchestrates BEE components

```
# BEE CONFIGURATION FILE #  
[DEFAULT]  
bee_workdir = $HOME/.config/beeflow  
  
[slurmrestd]  
slurm_socket = /tmp/slurm_${USER}_154.sock  
  
[graphdb]  
hostname = localhost  
dbpass = password  
bolt_port = 7741  
http_port = 7528  
https_port = 7527  
gdb_image = /usr/projects/beedev/neo4j-3-5-17.tar.gz  
gdb_image_mntdir = /tmp  
  
[workflow_manager]  
listen_port = 5054  
  
[task_manager]  
listen_port = 5104
```

BEEstart

- Manages configuration file complexity
- Launches BEE component servers

